

# Getting Started with



**BOSTON UNIVERSITY**  
**MACHINE INTELLIGENCE**  
**COMMUNITY**

Rachel Manzelli  
2/21/18

<https://goo.gl/8n4GbX>

# What is TensorFlow?

- Python framework focusing on data flow graphs, used for machine learning applications
- Builds static computation graphs
  - ◆ Supports graph-structured/tree-structured data
- Schedules operations automatically when run
- We will be using TensorFlow Core (higher level APIs are also available, i.e. `tf.contrib.learn`)

# Why TensorFlow?



# What You Need for Today

- A computer
- An installation of Python 2 or 3
- A text editor & means of compiling code

# Installing TensorFlow

If you have Python 2 or 3 installed, you can install TensorFlow via the command

```
>> pip install tensorflow
```

```
>> pip3 install tensorflow
```

in the command prompt.

You can also install via `virtualenv`.

(If these options cause issues, there is always the option of building from source...)



# Start programming in TensorFlow

Start programming in your favorite editor or IDE (Sublime, XCode, ... ).

TensorFlow is an environment in **Python**, so we will be using Python syntax.

To use TensorFlow, you will need to include

```
import tensorflow as tf
```

at the top of your .py file.

*But before we jump into that...*



# TensorFlow Concepts

- **Computational graph:** a series of operations arranged into a graph of nodes

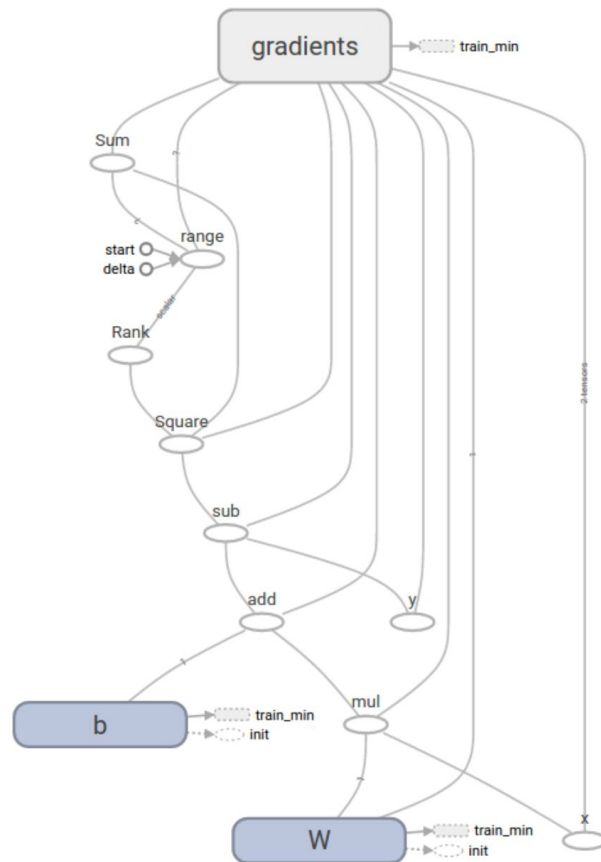
- ◆ **nodes**: take tensors as inputs and outputs a tensor
- ◆ **session** objects: `tf.Session.run()`

→ We create graphs in 2 steps:

1. Build (initialize)
2. Run (evaluate)

→ (Image taken from TensorBoard)

```
>> tensorboard --logdir=/your log dir
```





# TensorFlow Concepts

→ Tensors: representations of high-dimensional arrays

◆ **rank:** number of dimensions of a tensor




Figure 2: A 3rd-order tensor  $\underline{\mathbf{X}} \in \mathbb{R}^{I \times J \times K}$ , with entries  $x_{i,j,k}$  =



# How does TensorFlow manipulate data?

→ TensorFlow < 1.4

- ◆ **Placeholders (`tf.placeholder`)**: promise to provide a value later, but are initialized for the purpose of creating a graph
- ◆ **Feed dictionaries (`feed_dict`)**: allow for the node to be evaluated with multiple values
- ◆ **Variables (`tf.Variable`)**: trainable parameters, i.e. take in arbitrary values

→ TensorFlow 1.4 and 1.5

- ◆ Estimators and Datasets APIs



# Let's build a graph with TensorFlow!

- Open your editor again :)
- Import TensorFlow and initiate a session (using `tf.Session()`)
- Create a simple adder and multiplier by creating two nodes that will hold numbers, and nodes to store their output:

```
a = tf.placeholder(tf.float32)
b = tf.placeholder(tf.float32)
adder_node = a + b # + provides a shortcut for tf.add(a, b)
add_and_triple = adder_node * 3.
```



# Let's build a graph with TensorFlow!


- Let's compute some outputs using `tf.sesh.run()`.
- Evaluate these operations for multiple values using a feed dictionary:

```
print(sess.run(adder_node, {a: 3, b:4.5}))  
print(sess.run(adder_node, {a: [1,3], b: [2, 4]}))  
print(sess.run(add_and_triple, {a: 3, b:4.5}))
```

```
[7.5]  
[3. 7.]  
22.5
```



# What does our graph look like?




# Building a Trainable Graph

- Most applications in machine learning involve training, and updating parameters, which `tf.constant` cannot do. The values of constants in TensorFlow cannot change while the graph is evaluated.
- **`tf.Variable`** represents a trainable parameter in the graph (no concrete value)



# Building a Linear Model

- Let's start by building a very simple linear model.
- These models are based on a statistical process called **linear regression**: assuming the relationship between the data and parameters fits linearly.



# Building a Linear Model

- The only thing we know in this case is  $x$ . (We want to estimate  $y$ , but don't know the fit, which is represented by  $W$  and  $b$ ).
- So, we can create trainable parameters for the other parameters using **tf.Variable**:

```
W = tf.Variable([.3], dtype=tf.float32)
b = tf.Variable([-.3], dtype=tf.float32)
x = tf.placeholder(tf.float32)
linear_model = W * x + b
```





# Building a Linear Model

- When variables are created in TensorFlow, they require a special initializer.
- It is important to realize that we are not running the model (there are no parameters yet!) We just initialize the variables like this:

```
init = tf.global_variables_initializer()  
sess.run(init)
```

- Feed values into our model like before (x is the placeholder):

```
print(sess.run(linear_model, {x:[1,2,3,4]}))
```



# Evaluation of the Model

[0.            0.3            0.6            0.90000004]

- What does this mean? How can we evaluate the performance of this model?
- A common **loss function** for linear regression is **least squares**, which looks like this:

$$\sum_{i=1}^n (Y_i - \hat{Y}_i)^2.$$

... which squares the distance of the prediction from the ground truth value, and sums them all. We aim to minimize this sum (as close to zero as possible).

# Translating Loss to TensorFlow

- To evaluate the loss, we can translate our cost function to TensorFlow as follows (with  $y$  as a placeholder for the ground truth values):

```
y = tf.placeholder(tf.float32)
squared_deltas = tf.square(linear_model - y)
loss = tf.reduce_sum(squared_deltas)
print(sess.run(loss, {x:[1,2,3,4], y:[0,-1,-2,-3]}))
```

Ground truth values of  $y$   
for this linear model

- Our loss ends up being **23.66**



# Improving the Model

→ Manually assign the tensors  $W$  and  $b$  to -1 and 1, respectively.

```
W = tf.Variable([-1], dtype=tf.float32)
b = tf.Variable([1], dtype=tf.float32)
x = tf.placeholder(tf.float32)
linear_model = W * x + b
```

- These are the “perfect” values of our trainable parameters: the model will always predict  $y$  perfectly (with 0.0 loss) with these parameters (try it!)
- [https://github.com/pkmital/tensorflow\\_tutorials/blob/master/python/02\\_linear\\_regression.py](https://github.com/pkmital/tensorflow_tutorials/blob/master/python/02_linear_regression.py)



# Improving the Model

- Here, we know the values of  $W$  and  $b$  that will minimize the loss.
- The point of machine learning is for the model to figure out the perfect parameters through training and minimizing the cost function!

# A preview into how awesome TF is!

- Into your command shell of choice, clone TensorFlow's GitHub repository called "models":

```
>> git clone https://github.com/tensorflow/models
```

(This will take a few minutes - there's a lot in there to play around with!)

- For now, change into the `getting_started` directory:

```
>> cd models/samples/core/get_started
```



# Install dependencies

→ To run this example network, we need a few things:

```
>> pip install pandas: Data structure framework.
```

```
>> pip install keras: A TensorFlow wrapper.
```

```
>> pip install tensorflow --upgrade
```

Support for Keras was added in TensorFlow 1.4, so we need to make sure TensorFlow is updated.



# Before we run anything...

- What does this network do? Image classification!
  - ◆ Trains on images
  - ◆ Tests on images it hasn't seen before, and tries to classify the image as something it has learned
  - ◆ **Supervised learning problem**



# What does this network do?

- *Iris setosa*
- *Iris virginica*
- *Iris versicolor*



From left to right, *Iris setosa* (by [Radomil](#), CC BY-SA 3.0), *Iris versicolor* (by [Dlanglois](#), CC BY-SA 3.0), and *Iris virginica* (by [Frank Mayfield](#), CC BY-SA 2.0).

# Training and testing a network

- Into the command window, run the Python file `premade_estimator.py` in the current directory.

```
>> python premade_estimator.py
```

- This will train the network on the images for 1,000 iterations
- Then, it will test the network on test images
- Outputs loss and test accuracy by checking the predicted labels against the ground truth labels



```
INFO:tensorflow:Saving checkpoints for 1 into C:\Users\raman\AppData\Local\Temp\tmphrdx3y6c\model.ckpt.
INFO:tensorflow:loss = 279.39398, step = 1
INFO:tensorflow:global_step/sec: 852.418
INFO:tensorflow:loss = 16.906694, step = 101 (0.118 sec)
INFO:tensorflow:global_step/sec: 1061.02
INFO:tensorflow:loss = 10.27897, step = 201 (0.094 sec)
INFO:tensorflow:global_step/sec: 1108.16
INFO:tensorflow:loss = 9.171523, step = 301 (0.090 sec)
INFO:tensorflow:global_step/sec: 1108.16
INFO:tensorflow:loss = 7.0105577, step = 401 (0.090 sec)
INFO:tensorflow:global_step/sec: 1129.62
INFO:tensorflow:loss = 7.708337, step = 501 (0.089 sec)
INFO:tensorflow:global_step/sec: 1095.99
INFO:tensorflow:loss = 5.770026, step = 601 (0.091 sec)
INFO:tensorflow:global_step/sec: 1095.97
INFO:tensorflow:loss = 5.647806, step = 701 (0.091 sec)
INFO:tensorflow:global_step/sec: 1108.17
INFO:tensorflow:loss = 5.294636, step = 801 (0.090 sec)
INFO:tensorflow:global_step/sec: 1159.7
INFO:tensorflow:loss = 5.9896655, step = 901 (0.086 sec)
INFO:tensorflow:Saving checkpoints for 1000 into C:\Users\raman\AppData\Local\Temp\tmphrdx3y6c\model.ckpt.
INFO:tensorflow:Loss for final step: 3.7188513.
INFO:tensorflow:Starting evaluation at 2018-02-21-12:57:43
INFO:tensorflow:Restoring parameters from C:\Users\raman\AppData\Local\Temp\tmphrdx3y6c\model.ckpt-1000
INFO:tensorflow:Finished evaluation at 2018-02-21-12:57:43
INFO:tensorflow:Saving dict for global step 1000: accuracy = 0.96666664, average_loss = 0.058762528, global_step = 1000, loss = 1.7628758

Test set accuracy: 0.967

INFO:tensorflow:Restoring parameters from C:\Users\raman\AppData\Local\Temp\tmphrdx3y6c\model.ckpt-1000

Prediction is "Setosa" (99.6%), expected "Setosa"

Prediction is "Versicolor" (99.5%), expected "Versicolor"

Prediction is "Virginica" (97.3%), expected "Virginica"
PS C:\Users\raman\Documents\BUMIC\models\samples\core\get_started>
```

# Next time

## → **More concepts: Data and models!**

- ◆ How do we translate data into TensorFlow structures we saw today?
- ◆ What is a model? What types of models can we run with TensorFlow?

## → **More exercises/projects**

- ◆ Leave feedback and suggestions for what to do next time!



# Upcoming Events with BUMIC

- Impact of AI discussion group: 2/27/18
- Deep learning security lecture: 2/25/18
- Next TensorFlow session: 3/7/18 @ 7pm



Thank you to our sponsors!

