

Style Transfer with



BOSTON UNIVERSITY
MACHINE INTELLIGENCE
COMMUNITY

Rachel Manzelli
3/28/18

<https://goo.gl/bGiujV>

Last Time

- Learned the concepts of TensorFlow 1.2
- Built and evaluated simple graphs
- “Trained” a linear model

<https://github.com/bumic/TF-Workshops/tree/master/Workshop%201>



Today's Plan

- *Begin project*: implementation of neural style transfer in TensorFlow (will last 3 workshops)
- *Concepts*: convolutional neural networks, VGG-19 architecture, neural style transfer
- *Implementations*: structure of a neural network implementation, skeleton code for NST

What You Need For Today

- A computer & text editor
- Installations of Python 2 or 3, TensorFlow, numpy, scipy, and Pillow
- These can all be installed via `pip install`

```
>> pip install numpy
```

```
>> pip install scipy
```


```
>> pip install Pillow
```



A Brief Overview of Neural Networks

- ➔ Multilayer perceptrons: the “plain vanilla” of machine learning
- ➔ In a fully connected network, each neuron (a tensor in TensorFlow) in the next layer is connected to all neurons in the previous layer, and **feeds forward** information to “activate” the next neuron
- ➔ What is this information?






Neural Networks: Connections

- Weights: how “important” that neuron is
- Bias: how high the weighted sum needs to be for neuron to be meaningfully active
- Initialized randomly, then through training they are adjusted toward the correct value

$$Y = \sum (weight * input) + bias$$

$$\text{linear_model} = W * x + b$$





Inputs

Weights


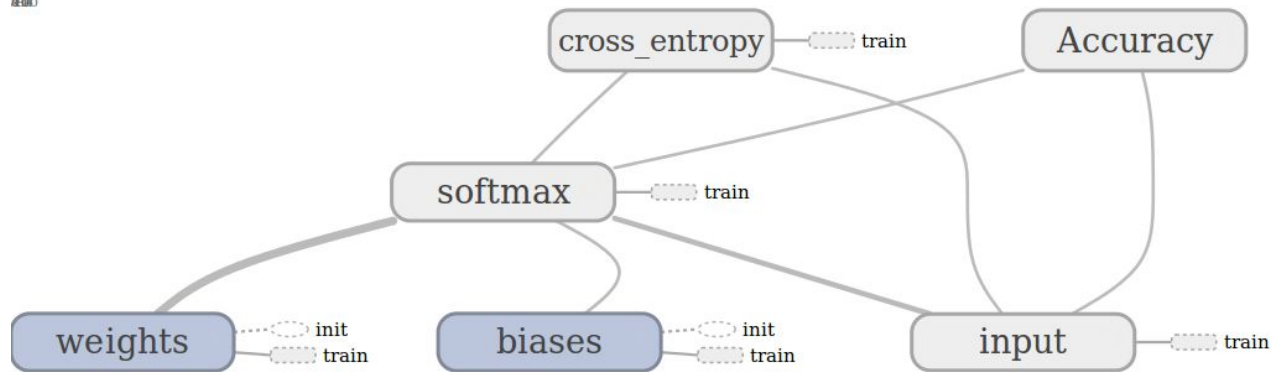
Summation and Bias

Activation

Output



28/08/20



M I C

Neural Networks: Loss functions

- Loss functions add a cost to the network's output, i.e. penalizing the network for performing poorly
 - ◆ Difference between output and desired output
 - ◆ We want to **minimize** this. There are many loss functions...

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

$$\mathcal{L}(y, \hat{y}) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$$

loss



Neural Networks: Training


- Training a network means minimizing this cost function more and more with each step.
- Gradient descent is called an optimizer: it allows us to adjust weights and biases by telling them to increase or decrease to minimize the cost function
- There are many other optimizers...

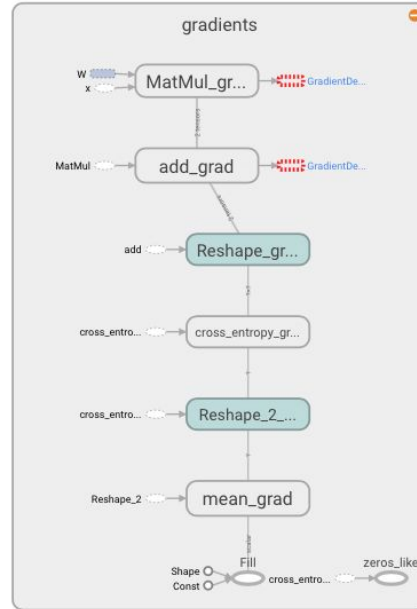

$$\mathbf{a}_{n+1} = \mathbf{a}_n - \gamma \nabla F(\mathbf{a}_n)$$

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t$$



Visual: Backpropagation







François Chollet ✓

@fchollet

Follow



"Neural networks" are a sad misnomer. They're neither neural nor even networks. They're chains of differentiable, parameterized geometric functions, trained with gradient descent (with gradients obtained via the chain rule). A small set of highschool-level ideas put together

11:58 AM - 12 Jan 2018

1,308 Retweets 3,413 Likes



116



1.3K



3.4K



M I C



François Chollet ✓

@fchollet | 81,895 followers

 Activity |  Map |  Analyse |  Visualise |  Alert

There's a point to be made that a neural network for food classification that can classify Tide pods as "not food" has already reached superhuman capabilities

 2 months ago · *Twitter Web Client* · en



M I C


Convolutional Neural Networks

- Neural networks with **convolutional layers**: use convolutions instead of matrix multiplications to produce output in hidden layers
- Each convolution layer takes an input image and produces a set of **feature maps** by **filtering**
 - ◆ Each filter is “convolved” across the width and height of the input volume, producing a 2-dimensional activation map of that filter.


Input layer

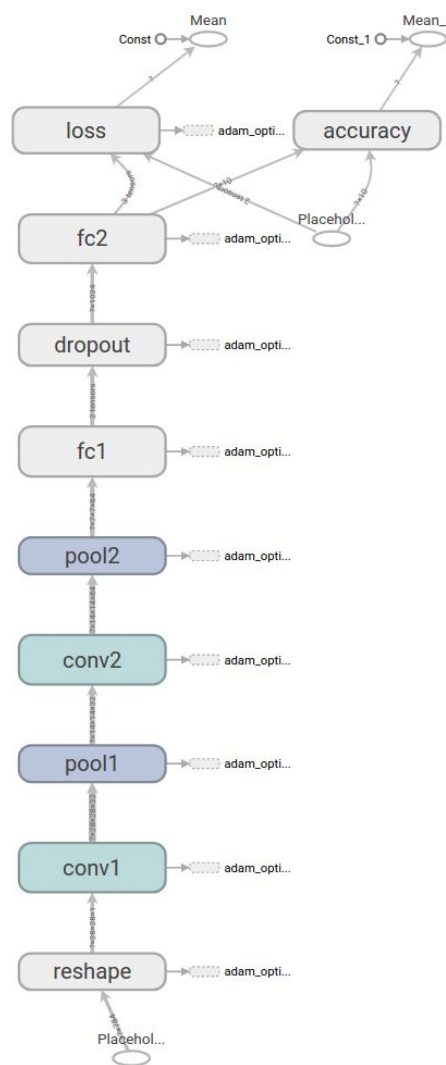


(S1) 4 feature maps



(C1) 4 feature maps







VGG-19

- A specific version of VGGNet with 19 weight layers
- VGGNet is a convolutional neural network originally used for image classification
 - ◆ uses 3×3 convolutional layers stacked on top of each other
 - ◆ reducing volume size is handled by max pooling
 - ◆ Two fully-connected layers, each with 4,096 nodes are followed by a softmax classifier.
- Original VGGNet paper:



<https://arxiv.org/abs/1409.1556>






Softmax

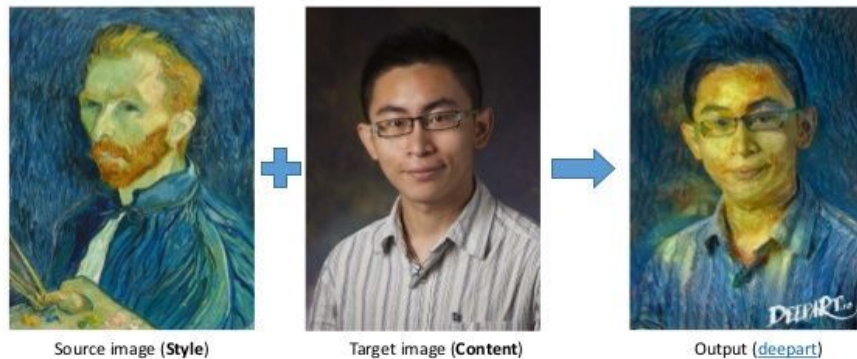
$$\sigma(x_j) = \frac{e^{x_j}}{\sum_i e^{x_i}}$$

Max pooling



What is Neural Style Transfer?

- Representations of style and content can be isolated and manipulated separately from a convolutional NN!
- Contains three main components:
 - ◆ Content image, style image, pastiche



A Neural Algorithm of Artistic Style [[Gatys et al. 2015](#)]






Content image:

Picture we want to transfer
style onto



Style image:

Artwork whose style we want to
transfer




Pastiche:

Final stylized image
(Initialized to random/white
noise)

Content image: loss and reconstruction

- **Content Loss** - Mean Squared Error between feature maps for a given convolutional layer l
 - P_{ij}^l : Feature map of original image
 - F_{ij}^l : Feature map of image content to be generated
- Gradient descent is used to minimize the content loss

$$L_{content}(\overrightarrow{p}, \overrightarrow{x}, l) = \frac{1}{2} \sum_{ij} (F_{ij}^l - P_{ij}^l)^2$$


Original image Generated image

Style image: loss function

- Style Loss

- Gram Matrix: inner product between two vectorised features maps i and j
- Computes correlations between different features

$$G_{ij}^l = \sum_k F_{ik}^l F_{jk}^l$$

- Minimize mean-squared distance between entries of Gram matrix of style image and generated image

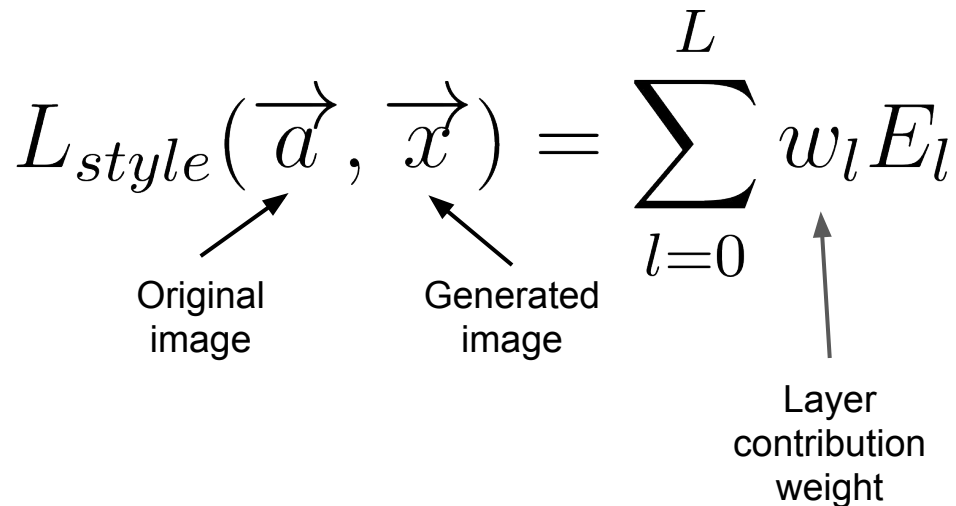
$$E_l = \frac{1}{4N_l^2 M_l^2} \sum_{i,j} (G_{ij}^l - A_{ij}^l)^2$$

Style image: reconstruction

- Total Style Loss across multiple layers
 - Weighted sum of styles losses across all layers of the convolutional network

$$L_{style}(\vec{a}, \vec{x}) = \sum_{l=0}^L w_l E_l$$

Original image Generated image Layer contribution weight

The diagram illustrates the formula for the total style loss. It features the equation $L_{style}(\vec{a}, \vec{x}) = \sum_{l=0}^L w_l E_l$. Three arrows point from text labels below to specific parts of the equation: one from 'Original image' to \vec{a} , one from 'Generated image' to \vec{x} , and one from 'Layer contribution weight' to w_l .

Total Loss Function for NST

- **Total Loss** - Jointly minimize error of a white noise image from
 - Content representation of content image \vec{p}
 - Style representation of style image \vec{a}

$$L_{total}(\vec{p}, \vec{a}, \vec{x}) = \alpha L_{content}(\vec{p}, \vec{x}) + \beta L_{style}(\vec{a}, \vec{x})$$

- α, β : weighting factors for content and style reconstruction respectively

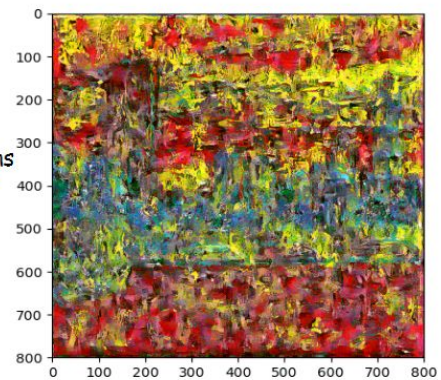
$$\mathbf{x}^* = \underset{\mathbf{x}}{\operatorname{argmin}} (\alpha \mathcal{L}_{\text{content}}(\mathbf{c}, \mathbf{x}) + \beta \mathcal{L}_{\text{style}}(\mathbf{s}, \mathbf{x}))$$



+



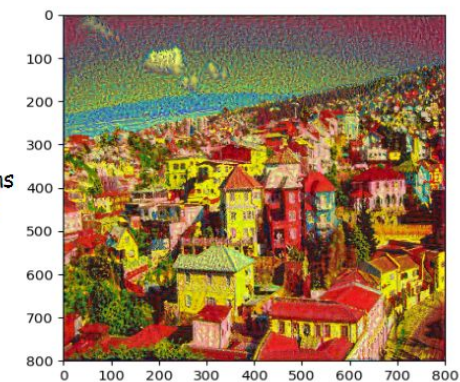
1000 iterations



+




1000 iterations



M I C

← Style emphasis α/β ratio Content emphasis →



VGG
content
layer

Implementation of Neural Style Transfer


- The implementation of neural style we will be using is located here:

<https://goo.gl/M1pBRx>

- Let's go to it and play around!



Implementation of Neural Style Transfer



Iteration 10




Iteration 300




Iteration 990


Implementation of Neural Style Transfer



Iteration 10



Iteration 300



Iteration 990

Let's (finally) start building it!

- Go to <https://github.com/anishathalye/neural-style>
- This is the NST implementation we will be rebuilding.
- Download the file at the bottom:

`imagenet-vgg-verydeep-19.mat`

- This is our pretrained VGGNet that we will be using in our implementation. (pre-training a convnet is time-consuming!)



The Pre-trained Network


- This network was trained on ImageNet (it can presumably classify images)
- Open the file in MATLAB
- If you don't have MATLAB, that's ok - Python to come!

Implementing a Neural Network

- *Architecting*: constructing the architecture of the neural network
- *Preprocessing*: preprocessing the data so that it is in a form we can train on
- *Training*: getting those perfect weights
- *Testing/generating*: using the trained network to generate a result

Implementation: Architecture

- Open a new file. This will be our initialization of the architecture of the network.
- First, we will initialize a dictionary of keys to refer back to the weight layers of our network, based on the .mat file.



Upcoming Events with BUMIC

- Impact of AI discussion group: 4/4/18
- Next TensorFlow session: 4/11/18 @ 7pm
 - ◆ Continuing NST: Finishing architecture and training the network



Thank you to our sponsors!

